1   Overview of Trust-region Methods

For nice figures, see [1].

   We just deal here with a small subset of trust-region methods, specifically approximating the cost function as quadratic using Newton's method, and using the Dogleg method and later to include Steihaug's method.

   The overall goal of a nonlinear optimization method is to iteratively find a local minimum of a nonlinear function

$$\hat{x} = \arg\min_x f(x)$$

where $f(x) \to \mathbb{R}$ is a scalar function. In GTSAM, the variables x could be manifold or Lie group elements, so in this document we only work with *increments* $\delta x \in \mathbb{R}^n$ in the tangent space. In this document we specifically deal with *trust-region* methods, which at every iteration attempt to find a good increment $\|\delta x\| \leq \Delta$ within the "trust radius" Δ.

   Further, most nonlinear optimization methods, including trust region methods, deal with an approximate problem at every iteration. Although there are other choices (such as quasi-Newton), the Newton's method approximation is, given an estimate $x^{(k)}$ of the variables x,

$$f\left(x^{(k)} \oplus \delta x\right) \approx M^{(k)}(\delta x) = f^{(k)} + g^{(k)\mathsf{T}}\delta x + \frac{1}{2}\delta x^\mathsf{T} G^{(k)}\delta x, \tag{1}$$

where $f^{(k)} = f\left(x^{(k)}\right)$ is the function at $x^{(k)}$, $g^{(x)} = \left.\frac{\partial f}{\partial x}\right|_{x^{(k)}}$ is its gradient, and $G^{(k)} = \left.\frac{\partial^2 f}{\partial x^2}\right|_{x^{(k)}}$ is its Hessian (or an approximation of the Hessian).

   Trust-region methods adaptively adjust the trust radius Δ so that within it, M is a good approximation of f, and then never step beyond the trust radius in each iteration. When the true minimum is within the trust region, they converge quadratically like Newton's method. At each iteration k, they solve the *trust-region subproblem* to find a proposed update δx inside the trust radius Δ, which decreases the approximate function $M^{(k)}$ as much as possible. The proposed update is only accepted if the true function f decreases as well. If the decrease of M matches the decrease of f well, the size of the trust region is increased, while if the match is not close the trust region size is decreased.

   Minimizing Eq. 1 is itself a nonlinear optimization problem, so there are various methods for approximating it, including Dogleg and Steihaug's method.

2   Adapting the Trust Region Size

As mentioned in the previous section, we increase the trust region size if the decrease in the model function M matches the decrease in the true cost function S very closely, and decrease it if they do not match closely. The closeness of this match is measured with the *gain ratio*,

$$\rho = \frac{f(x) - f(x \oplus \delta x_d)}{M(0) - M(\delta x_d)},$$

where δx_d is the proposed dogleg step to be introduced next. The decrease in the model function is always non-negative, and as the decrease in f approaches it, ρ approaches 1. If the true cost function increases, ρ will be negative, and if the true cost function decreases even more than predicted by M, then ρ will be greater than 1. A typical update rule, as per Lec. 7-1.2 of [1] is:

$$\Delta_{k+1} \leftarrow \begin{cases} \Delta_k/4 & \rho < 0.25 \\ \min\left(2\Delta_k, \Delta_{max}\right), & \rho > 0.75 \\ \Delta_k & 0.75 > \rho > 0.25 \end{cases}$$

where $\Delta_k \triangleq \|\delta x_d\|$. Note that the rule is designed to ensure that Δ_k never exceeds the maximum trust region size Δ_{max}.

3 Dogleg

Dogleg minimizes an approximation of Eq. 1 by considering three possibilities using two points - the minimizer $\delta x_u^{(k)}$ of $M^{(k)}$ along the negative gradient direction $-g^{(k)}$, and the overall Newton's method minimizer $\delta x_n^{(k)}$ of $M^{(k)}$. When the Hessian $G^{(k)}$ is positive, the magnitude of $\delta x_u^{(k)}$ is always less than that of $\delta x_n^{(k)}$, meaning that the Newton's method step is "more adventurous". How much we step towards the Newton's method point depends on the trust region size:

1. If the trust region is smaller than $\delta x_u^{(k)}$, we step in the negative gradient direction but only by the trust radius.

2. If the trust region boundary is between $\delta x_u^{(k)}$ and $\delta x_n^{(k)}$, we step to the linearly-interpolated point between these two points that intersects the trust region boundary.

3. If the trust region boundary is larger than $\delta x_n^{(k)}$, we step to $\delta x_n^{(k)}$.

To find the intersection of the line between $\delta x_u^{(k)}$ and $\delta x_n^{(k)}$ with the trust region boundary, we solve a quadratic roots problem,

$$\Delta = \|(1 - \tau)\,\delta x_u + \tau \delta x_n\|$$
$$\Delta^2 = (1 - \tau)^2\,\delta x_u^\mathsf{T}\delta x_u + 2\tau\,(1 - \tau)\,\delta x_u^\mathsf{T}\delta x_n + \tau^2\delta x_n^\mathsf{T}\delta x_n$$
$$0 = uu - 2\tau uu + \tau^2 uu + 2\tau un - 2\tau^2 un + \tau^2 nn - \Delta^2$$
$$0 = (uu - 2un + nn)\,\tau^2 + (2un - 2uu)\,\tau - \Delta^2 + uu$$
$$\tau = \frac{-(2un - 2uu) \pm \sqrt{(2un - 2uu)^2 - 4\,(uu - 2un + nn)\,(uu - \Delta^2)}}{2\,(uu - un + nn)}$$

From this we take whichever possibility for τ such that $0 < \tau < 1$.

To find the steepest-descent minimizer $\delta x_u^{(k)}$, we perform line search in the gradient direction on the approximate function M,

$$\delta x_u^{(k)} = \frac{-g^{(k)\mathsf{T}} g^{(k)}}{g^{(k)\mathsf{T}} G^{(k)} g^{(k)}} g^{(k)} \tag{2}$$

Thus, mathematically, we can write the dogleg update $\delta x_d^{(k)}$ as

$$\delta x_d^{(k)} = \begin{cases} -\frac{\Delta}{\left\| \delta x_u^{(k)} \right\|} \delta x_u^{(k)}, & \Delta < \left\| \delta x_u^{(k)} \right\| \\ \left(1 - \tau^{(k)}\right) \delta x_u^{(k)} + \tau^{(k)} \delta x_n^{(k)}, & \left\| \delta x_u^{(k)} \right\| < \Delta < \left\| \delta x_n^{(k)} \right\| \\ \delta x_n^{(k)}, & \left\| \delta x_n^{(k)} \right\| < \Delta \end{cases}$$

4   Working with M as a Bayes' Net

When we have already eliminated a factor graph into a Bayes' Net, we have the same quadratic error function $M^{(k)}(\delta x)$, but it is in a different form:

$$M^{(k)}(\delta x) = \frac{1}{2} \left\| Rx - d \right\|^2,$$

where R is an upper-triangular matrix (stored as a set of sparse block Gaussian conditionals in GTSAM), and d is the r.h.s. vector. The gradient and Hessian of M are then

$$g^{(k)} = R^\mathsf{T} (Rx - d)$$
$$G^{(k)} = R^\mathsf{T} R$$

In GTSAM, because the Bayes' Net is not dense, we evaluate Eq. 2 in an efficient way. Rewriting the denominator (leaving out the (k) superscript) as

$$g^\mathsf{T} G g = \sum_i (R_i g)^\mathsf{T} (R_i g),$$

where i indexes over the conditionals in the Bayes' Net (corresponding to blocks of rows of R) exploits the sparse structure of the Bayes' Net, because it is easy to only include the variables involved in each i^{th} conditional when multiplying them by the corresponding elements of g.

References

[1] Raphael Hauser. Lecture notes on unconstrained optimization. link, 2006.